

Part 2 — Solutions for Lecture 4

TECH2: Introduction to Programming, Data, and Information Technology

Richard Foltyn

Norwegian School of Economics (NHH)

October 23, 2024

See GitHub repository for notebooks and data:

<https://github.com/richardfoltyn/TECH2-H24>

Contents

1 Concatenating and merging data	1
1.1 Concatenation	1
1.2 Merging and joining data sets	4
2 Dealing with missing values	7

1 Concatenating and merging data

1.1 Concatenation

- Your turn.**
1. Create a new Series with observations ['C1', 'C2'].
 2. Using the previously created Series a and b, concatenate all three objects along the row axis and create a new (unique) index.
 3. Repeat the previous step, but now concatenate along the column axis. Assign the column names 'Column1', 'Column2', and 'Column3'.

Solution.

```
[2]: import pandas as pd

# Recreate Series a, b:
# Create first series of 3 observations
a = pd.Series(['A1', 'A2', 'A3'])
# Create second series with 5 observations
b = pd.Series([f'B{i}' for i in range(5)])

[3]: # Create Series c
c = pd.Series(['C1', 'C2'])

[4]: # Concatenate Series a, b, c and reset the index
s = pd.concat((a, b, c)).reset_index(drop=True)
s
```

```
[4]: 0    A1
      1    A2
      2    A3
      3    B0
      4    B1
      5    B2
      6    B3
      7    B4
      8    C1
      9    C2
      dtype: object
```

```
[5]: s = pd.concat((a, b, c), axis=1, keys=['Column1', 'Column2', 'Column3'])
      s
```

```
[5]:   Column1 Column2 Column3
0      A1      B0      C1
1      A2      B1      C2
2      A3      B2      NaN
3      NaN      B3      NaN
4      NaN      B4      NaN
```

Your turn. Use the data files located in the folder ../data/FRED to perform the following tasks:

1. Load the data in FRED_monthly_1950.csv and FRED_monthly_1960.csv into two different DataFrames. The files contain monthly macroeconomic time series for the 1950s and 1960s, respectively.

Hint: Use `pd.read_csv(..., parse_dates=['DATE'])` to automatically parse strings stored in the DATE column as dates.

2. Concatenate these DataFrames along the row dimension to get a total of 240 observations.
3. Set the column DATE as index for the newly created DataFrame.

Solution.

Part (1)

```
[83]: # Path to data folder
      DATA_PATH = '/home/richard/repos/teaching/TECH2-H24/data/FRED'
```

```
[84]: import pandas as pd

      # Load data from the 1950s
      df1 = pd.read_csv(f'{DATA_PATH}/FRED_monthly_1950.csv', parse_dates=['DATE'])
      df1.head(5)
```

```
[84]:   DATE    CPI  UNRATE  FEDFUNDS  REALRATE  LFPART
0 1950-01-01  23.5     6.5      NaN      NaN     58.9
1 1950-02-01  23.6     6.4      NaN      NaN     58.9
2 1950-03-01  23.6     6.3      NaN      NaN     58.8
3 1950-04-01  23.6     5.8      NaN      NaN     59.2
4 1950-05-01  23.8     5.5      NaN      NaN     59.1
```

```
[85]: # Load data from the 1960s
      df2 = pd.read_csv(f'{DATA_PATH}/FRED_monthly_1960.csv', parse_dates=['DATE'])
      df2.head(5)
```

```
[85]:
```

	DATE	CPI	UNRATE	FEDFUNDS	REALRATE	LFPART
0	1960-01-01	29.4	5.2	4.0	NaN	59.1
1	1960-02-01	29.4	4.8	4.0	NaN	59.1
2	1960-03-01	29.4	5.4	3.8	NaN	58.5
3	1960-04-01	29.5	5.2	3.9	NaN	59.5
4	1960-05-01	29.6	5.1	3.8	NaN	59.5

Part (2)

```
[86]: # Concatenate data sets along the first dimension (rows)
df = pd.concat((df1, df2), axis=0)
```

```
[87]: # First half contains data from the 1950s
df.head(5)
```

```
[87]:
```

	DATE	CPI	UNRATE	FEDFUNDS	REALRATE	LFPART
0	1950-01-01	23.5	6.5	NaN	NaN	58.9
1	1950-02-01	23.6	6.4	NaN	NaN	58.9
2	1950-03-01	23.6	6.3	NaN	NaN	58.8
3	1950-04-01	23.6	5.8	NaN	NaN	59.2
4	1950-05-01	23.8	5.5	NaN	NaN	59.1

```
[88]: # Second half contains data from the 1960s
df.tail(5)
```

```
[88]:
```

	DATE	CPI	UNRATE	FEDFUNDS	REALRATE	LFPART
115	1969-08-01	36.9	3.5	9.2	NaN	60.3
116	1969-09-01	37.1	3.7	9.2	NaN	60.3
117	1969-10-01	37.3	3.7	9.0	NaN	60.4
118	1969-11-01	37.5	3.5	8.8	NaN	60.2
119	1969-12-01	37.7	3.5	9.0	NaN	60.2

Part (3)

Note that the index of the newly created DataFrame is not unique:

```
[89]: # Select rows at index 0: returns 2 (!) different rows
df.loc[0]
```

```
[89]:
```

	DATE	CPI	UNRATE	FEDFUNDS	REALRATE	LFPART
0	1950-01-01	23.5	6.5	NaN	NaN	58.9
0	1960-01-01	29.4	5.2	4.0	NaN	59.1

```
[90]: # Set Date as new (unique!) index
df = df.set_index('DATE')
df.head(10)
```

```
[90]:
```

	CPI	UNRATE	FEDFUNDS	REALRATE	LFPART
DATE					
1950-01-01	23.5	6.5	NaN	NaN	58.9
1950-02-01	23.6	6.4	NaN	NaN	58.9
1950-03-01	23.6	6.3	NaN	NaN	58.8
1950-04-01	23.6	5.8	NaN	NaN	59.2
1950-05-01	23.8	5.5	NaN	NaN	59.1
1950-06-01	23.9	5.4	NaN	NaN	59.4
1950-07-01	24.1	5.0	NaN	NaN	59.1
1950-08-01	24.2	4.5	NaN	NaN	59.5
1950-09-01	24.3	4.4	NaN	NaN	59.2
1950-10-01	24.5	4.2	NaN	NaN	59.4

1.2 Merging and joining data sets

Your turn. Use the data files located in the folder `../data/FRED` to perform the following tasks:

1. Load the data in `CPI.csv` and `GDP.csv` into two different DataFrames. The files contain monthly data for the Consumer Price Index (CPI) and quarterly data for GDP, respectively.

Hint: Use `pd.read_csv(..., parse_dates=['DATE'])` to automatically parse strings stored in the `DATE` column as dates.

2. Merge the CPI with the GDP time series with `merge()` using a left join (`how='left'`). How many observations does the resulting DataFrame have?
3. Merge the CPI with the GDP time series with `merge()` using an inner join (`how='inner'`). How many observations does the resulting DataFrame have, and why is this different from the previous case?

Solution.

Part (1)

```
[91]: # Path to data folder
DATA_PATH = '/home/richard/repos/teaching/TECH2-H24/data/FRED'
```

```
[92]: import pandas as pd

cpi = pd.read_csv(f'{DATA_PATH}/CPI.csv', parse_dates=['DATE'])
cpi.head(5)
```

```
[92]:      DATE  CPI
0  1947-01-01  21.5
1  1947-02-01  21.6
2  1947-03-01  22.0
3  1947-04-01  22.0
4  1947-05-01  22.0
```

```
[93]: gdp = pd.read_csv(f'{DATA_PATH}/GDP.csv', parse_dates=['DATE'])
gdp.head(5)
```

```
[93]:      DATE      GDP
0  1947-01-01  2182.7
1  1947-04-01  2176.9
2  1947-07-01  2172.4
3  1947-10-01  2206.5
4  1948-01-01  2239.7
```

Part (2)

```
[94]: # Merge so that left DataFrame determines resulting index
df = pd.merge(cpi, gdp, on='DATE', how='left')
df.head(12)
```

```
[94]:      DATE  CPI      GDP
0  1947-01-01  21.5  2182.7
```

1	1947-02-01	21.6	NaN
2	1947-03-01	22.0	NaN
3	1947-04-01	22.0	2176.9
4	1947-05-01	22.0	NaN
5	1947-06-01	22.1	NaN
6	1947-07-01	22.2	2172.4
7	1947-08-01	22.4	NaN
8	1947-09-01	22.8	NaN
9	1947-10-01	22.9	2206.5
10	1947-11-01	23.1	NaN
11	1947-12-01	23.4	NaN

```
[95]: # Number of observations
N = len(df)
print(f'Number of observations with left join: {N:,d}')
```

Number of observations with left join: 932

Part (3)

```
[96]: # Drop columns with missing observations in GDP
df = pd.merge(cpi, gdp, on='DATE', how='inner')
df.head(12)
```

```
[96]:
```

	DATE	CPI	GDP
0	1947-01-01	21.5	2182.7
1	1947-04-01	22.0	2176.9
2	1947-07-01	22.2	2172.4
3	1947-10-01	22.9	2206.5
4	1948-01-01	23.7	2239.7
5	1948-04-01	23.8	2276.7
6	1948-07-01	24.4	2289.8
7	1948-10-01	24.3	2292.4
8	1949-01-01	24.0	2260.8
9	1949-04-01	23.9	2253.1
10	1949-07-01	23.7	2276.4
11	1949-10-01	23.7	2257.4

```
[97]: # Number of observations
N = len(df)
print(f'Number of observations with inner join: {N:,d}')
```

Number of observations with inner join: 310

The inner join drops all dates from `cpi` which are not present in the `gdp` DataFrame, hence the number of rows in the merged DataFrame is only a third of the original data (since the GDP data is quarterly).

Your turn. Use the data files located in the folder `../data/FRED` to perform the following tasks:

1. Load the data in `CPI.csv` and `GDP.csv` into two different DataFrames. The files contain monthly data for the Consumer Price Index (CPI) and quarterly data for GDP, respectively.

Hint: Use `pd.read_csv(..., parse_dates=['DATE'])` to automatically parse strings stored in the `DATE` column as dates.

2. Set the `DATE` column as the index for each of the two DataFrames.
3. Merge the CPI with the GDP time series with `join()`. Do this with both a left and an inner join.

Solution.

Part (1)

```
[98]: # Path to data folder
DATA_PATH = '/home/richard/repos/teaching/TECH2-H24/data/FRED'
```

```
[99]: import pandas as pd

cpi = pd.read_csv(f'{DATA_PATH}/CPI.csv', parse_dates=['DATE'])
# Alternatively, we can set the index directly when loading the data
# cpi = pd.read_csv(f'{DATA_PATH}/CPI.csv', parse_dates=['DATE'], index_col='DATE')
cpi.head(5)
```

```
[99]:      DATE    CPI
0 1947-01-01  21.5
1 1947-02-01  21.6
2 1947-03-01  22.0
3 1947-04-01  22.0
4 1947-05-01  22.0
```

```
[100]: gdp = pd.read_csv(f'{DATA_PATH}/GDP.csv', parse_dates=['DATE'])
# Alternatively, we can set the index directly when loading the data
# gdp = pd.read_csv(f'{DATA_PATH}/GDP.csv', parse_dates=['DATE'], index_col='DATE')
gdp.head(5)
```

```
[100]:      DATE    GDP
0 1947-01-01 2182.7
1 1947-04-01 2176.9
2 1947-07-01 2172.4
3 1947-10-01 2206.5
4 1948-01-01 2239.7
```

Part (2)

If we didn't specify the index columns using `index_col` as an argument to `pd.read_csv()`, we can set the index after loading the data.

```
[101]: # Set DATE column as index
cpi = cpi.set_index('DATE')
gdp = gdp.set_index('DATE')
```

Part (3)

```
[102]: # Perform left join (the default)
df = cpi.join(gdp)
df.head(10)
```

```
[102]:      CPI    GDP
DATE
1947-01-01  21.5 2182.7
1947-02-01  21.6   NaN
1947-03-01  22.0   NaN
1947-04-01  22.0 2176.9
1947-05-01  22.0   NaN
1947-06-01  22.1   NaN
1947-07-01  22.2 2172.4
```

```
1947-08-01  22.4      NaN
1947-09-01  22.8      NaN
1947-10-01  22.9  2206.5
```

```
[103]: # Perform inner join
df = cpi.join(gdp, how='inner')
df.head(10)
```

```
[103]:      CPI      GDP
DATE
1947-01-01  21.5  2182.7
1947-04-01  22.0  2176.9
1947-07-01  22.2  2172.4
1947-10-01  22.9  2206.5
1948-01-01  23.7  2239.7
1948-04-01  23.8  2276.7
1948-07-01  24.4  2289.8
1948-10-01  24.3  2292.4
1949-01-01  24.0  2260.8
1949-04-01  23.9  2253.1
```

2 Dealing with missing values

Your turn. Use the data files located in the folder ../data/FRED to perform the following tasks:

1. Load the data in CPI.csv and GDP.csv into two different DataFrames. The files contain monthly data for the Consumer Price Index (CPI) and quarterly data for GDP, respectively.

Hint: Use `pd.read_csv(..., parse_dates=['DATE'])` to automatically parse strings stored in the DATE column as dates.

2. Merge the CPI with the GDP time series with `merge()` using a left join. This creates missing values in the GDP column.
3. Impute the missing GDP values using `interpolate()` and replace the missing values in column GDP.

Solution.

Part (1)

```
[2]: # Path to data folder
DATA_PATH = '/home/richard/repos/teaching/TECH2-H24/data/FRED'
```

```
[3]: import pandas as pd

# Load CPI data
cpi = pd.read_csv(f'{DATA_PATH}/CPI.csv', parse_dates=['DATE'])

# Load GDP data
gdp = pd.read_csv(f'{DATA_PATH}/GDP.csv', parse_dates=['DATE'])
```

Part (2)

```
[7]: # Merge CPI and GDP into a single DataFrame, use keys from CPI
df = pd.merge(cpi, gdp, how='left')

# Print first 12 months
df.head(12)
```

```
[7]:
```

	DATE	CPI	GDP
0	1947-01-01	21.5	2182.7
1	1947-02-01	21.6	NaN
2	1947-03-01	22.0	NaN
3	1947-04-01	22.0	2176.9
4	1947-05-01	22.0	NaN
5	1947-06-01	22.1	NaN
6	1947-07-01	22.2	2172.4
7	1947-08-01	22.4	NaN
8	1947-09-01	22.8	NaN
9	1947-10-01	22.9	2206.5
10	1947-11-01	23.1	NaN
11	1947-12-01	23.4	NaN

Since GDP data is available on quarterly frequency, only every third month contains non-missing values.

Part (3)

```
[12]: # Linearly interpolate missing value
df['GDP'] = df['GDP'].interpolate(method='linear')

# Print first 12 months to confirm that missing values are gone
df.head(12)
```

```
[12]:
```

	DATE	CPI	GDP
0	1947-01-01	21.5	2182.700000
1	1947-02-01	21.6	2180.766667
2	1947-03-01	22.0	2178.833333
3	1947-04-01	22.0	2176.900000
4	1947-05-01	22.0	2175.400000
5	1947-06-01	22.1	2173.900000
6	1947-07-01	22.2	2172.400000
7	1947-08-01	22.4	2183.766667
8	1947-09-01	22.8	2195.133333
9	1947-10-01	22.9	2206.500000
10	1947-11-01	23.1	2217.566667
11	1947-12-01	23.4	2228.633333
